

# Analysis of intrinsic vowel duration in Northwestern Italian

Stefano Coretta

2023-12-13

## Attach packages

```
library(tidyverse)
theme_set(theme_light())
library(magrittr)
library(coretta2018itaegg)
library(brms)
library(posterior)
library(tidybayes)
library(marginaleffects)
library(ggdist)
library(mgcv)
library(tidygam)

my_seed <- 9899
```

## Read data

```
data("formants")

formants %<>% mutate(
  duration = duration * 1000,
  vowel = as.factor(label),
  duration_z = as.vector(scale(duration)),
  duration_log = log(duration),
  duration_logz = as.vector(scale(log(duration))),
```

```

f13_z = as.vector(scale(f13)),
f23_z = as.vector(scale(f23)),
speaker = as.factor(speaker)
)
contrasts(formants$vowel) <- "contr.sum"

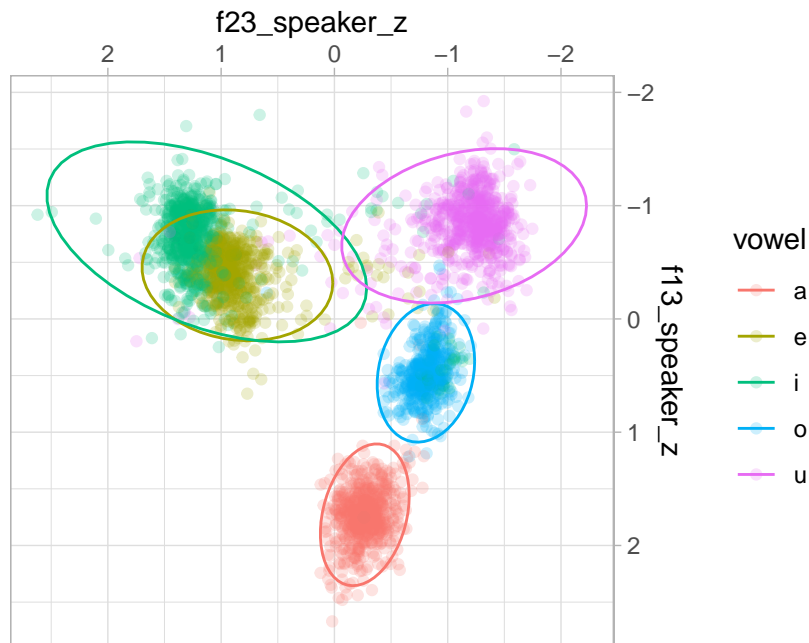
```

## Plotting

```

formants %>%
  group_by(speaker) %>%
  mutate(
    f13_speaker_z = as.vector(scale(f13)),
    f23_speaker_z = as.vector(scale(f23))
  ) %>%
  ggplot(aes(f23_speaker_z, f13_speaker_z, colour = vowel)) +
  geom_point(alpha = 0.2) +
  stat_ellipse(type = "norm") +
  scale_x_reverse(position = "top") + scale_y_reverse(position = "right") +
  coord_fixed()

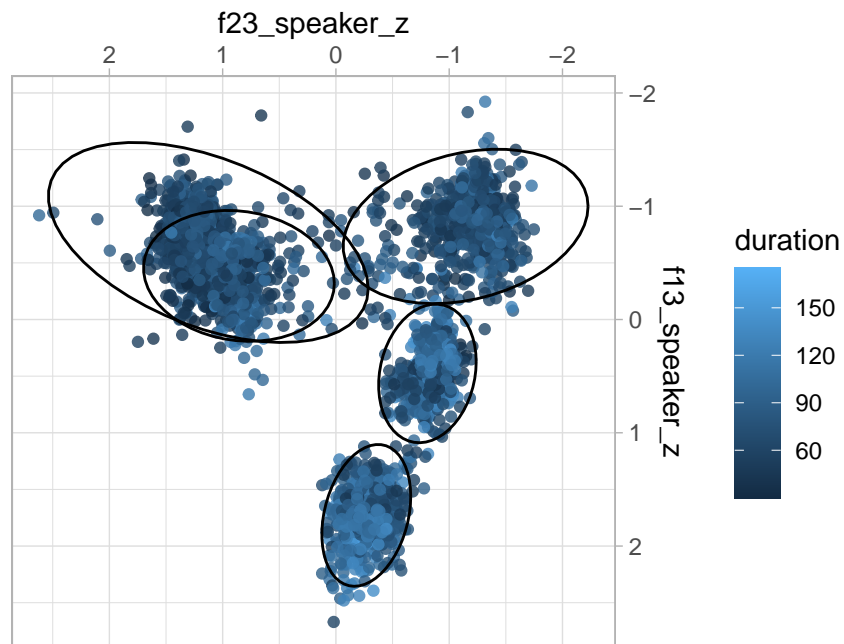
```



```

formants %>%
  group_by(speaker) %>%
  mutate(
    f13_speaker_z = as.vector(scale(f13)),
    f23_speaker_z = as.vector(scale(f23))
  ) %>%
  ggplot(aes(f23_speaker_z, f13_speaker_z)) +
  geom_point(aes(colour = duration), alpha = 0.8) +
  stat_ellipse(aes(group = vowel), type = "norm") +
  scale_x_reverse(position = "top") + scale_y_reverse(position = "right") +
  coord_fixed()


```




```

formants %>%
  ggplot(aes(f13_z, duration_z)) +
  geom_density_2d_filled()

```



```
formants %>%
  ggplot(aes(f13_z, duration_z, colour = vowel)) +
  geom_point(alpha = 0.2) +
  geom_smooth(method = "lm", formula = y ~ x)
```



## Linear modelling

### Prior predictive checks


The outcome `duration_logz` and predictor `f13_z` are z-scored and `vowel` is sum coded so that Intercept is the grand mean.

I am using relatively weakly informative priors.


```
priors <- c(
  prior(normal(0, 1), class = Intercept),
  prior(normal(0, 1), class = b),
  prior(cauchy(0, 0.1), class = sigma),
  prior(lkj(2), class = cor),
  prior(cauchy(0, 0.1), class = sd)
)

bm_1_priors <- brm(
  duration_logz ~
    vowel * f13_z +
    (vowel * f13_z | speaker),
  family = gaussian,
  data = formants,
  prior = priors,
  cores = 4,
  threads = threading(2),
  backend = "cmdstanr",
  sample_prior = "only",
  file = "data/cache/bm_1_priors",
)

conditional_effects(bm_1_priors, "f13_z:vowel")
```



```
conditional_effects(bm_1_priors, "vowel")
```



## Model fit

```
bm_1 <- brm(
  duration_logz ~
    vowel * f13_z +
    (vowel * f13_z | speaker),
  family = gaussian,
  data = formants,
  prior = priors,
  cores = 4,
  threads = threading(2),
  backend = "cmdstanr",
  file = "data/cache/bm_1",
)

fixef(bm_1, probs = c(0.05, 0.95))
```

	Estimate	Est.Error	Q5	Q95
Intercept	-0.15456092	0.14810441	-0.39139975	0.09485949
vowel1	-0.28607100	0.13267976	-0.51491880	-0.07680009
vowel2	0.35434327	0.05833619	0.25963235	0.45007120
vowel3	-0.32030985	0.05730238	-0.41276840	-0.22539825
vowel4	0.47457028	0.07210259	0.35954355	0.59458675
f13_z	0.36700826	0.05277162	0.28164785	0.45174305
vowel1:f13_z	0.35767128	0.08984698	0.21301775	0.50970635
vowel2:f13_z	0.12961894	0.09920732	-0.03739698	0.28873440
vowel3:f13_z	-0.05189318	0.06790845	-0.16322630	0.06080566
vowel4:f13_z	-0.01187440	0.10306594	-0.18354195	0.15770550

Let's get the estimated effect of vowel quality for /u/.

```
bm_1_draws <- as_draws_df(bm_1) %>%
  mutate(
    vowel5 = b_Intercept - b_vowel1 - b_vowel2 - b_vowel3 - b_vowel4,
    b_vowel5 = vowel5 - b_Intercept
  )

quantile(bm_1_draws$b_vowel5, probs = c(0.05, 0.95))
```

5%	95%
-0.37541152	-0.07203285

```

bm_1 %>%
  as_draws_df() %>%
  select(b_vowel1:b_vowel4) %>%
  pivot_longer(b_vowel1:b_vowel4) %>%
  group_by(name) %>%
  summarise(
    cri95 = list(round(quantile2(value, probs = c(0.025, 0.975)), 2)),
    cri90 = list(round(quantile2(value, probs = c(0.05, 0.95)), 2)),
    cri80 = list(round(quantile2(value, probs = c(0.1, 0.9)), 2)),
    cri60 = list(round(quantile2(value, probs = c(0.2, 0.8)), 2))
  ) %>%
  knitr::kable(format = "latex") %>% cat(sep = "\n")

```

Warning: Dropping 'draws\_df' class as required metadata was removed.

```

\begin{tabular}{l|l|l|l|l|l}
\hline
name & cri95 & cri90 & cri80 & cri60 & \\
\hline
b\_vowel1 & -0.55, -0.04 & -0.51, -0.08 & -0.46, -0.12 & -0.40, -0.17 & \\
\hline
b\_vowel2 & 0.24, 0.47 & 0.26, 0.45 & 0.28, 0.43 & 0.3, 0.4 & \\
\hline
b\_vowel3 & -0.43, -0.21 & -0.41, -0.23 & -0.39, -0.25 & -0.37, -0.27 & \\
\hline
b\_vowel4 & 0.33, 0.62 & 0.36, 0.59 & 0.38, 0.57 & 0.42, 0.53 & \\
\hline
\end{tabular}

```


## Model plotting

```


conditional_effects(bm_1, "f13_z:vowel", prob = 0.9)

```





```
conditional_effects(bm_1, "f13_z:vowel", spaghetti = TRUE, ndraws = 100, prob = 0.9)
```



We need to get the predicted draws to convert duration and F1 back to ms and hz. Note that

duration was logged then scaled.

```
seq_minmax <- function(x, by = 1) {  
  seq(min(x), max(x), by = by)  
}  
  
bm_1_grid <- expand_grid(  
  vowel = levels(formants$vowel),  
  f13_z = seq_minmax(formants$f13_z, 0.5)  
)  
  
bm_1_preds <- epred_draws(bm_1, newdata = bm_1_grid, re_formula = NA) %>%  
  mutate(  
    duration_log = .epred * sd(formants$duration_log) + mean(formants$duration_log),  
    duration = exp(duration_log),  
    f13 = f13_z * sd(formants$f13) + mean(formants$f13)  
  )
```


Let's also calculate the mean F1 values for each vowel, to be added in the plot below.

```
vmean_f13 <- formants %>%  
  group_by(vowel) %>%  
  summarise(f13_mean = mean(f13))  
vmean_f13z <- formants %>%  
  group_by(vowel) %>%  
  summarise(f13z_mean = mean(f13_z))
```


We can now plot the model predictions in the original scale.

```
bm_1_preds %>%  
  group_by(vowel, f13) %>%  
  summarise(  
    duration_ms = median(duration),  
    # Get the 90% CrI  
    q0.05 = quantile(duration, probs = 0.05),  
    q0.90 = quantile(duration, probs = 0.95),  
    .groups = "drop"  
  ) %>%  
  ggplot(aes(f13, duration_ms)) +  
  geom_ribbon(aes(ymin = q0.05, ymax = q0.90, fill = vowel), alpha = 0.25) +  
  geom_line(aes(colour = vowel), linewidth = 1) +  
  facet_wrap(~vowel) +
```


```
scale_colour_viridis_d() +
scale_fill_viridis_d()
```



```
bm_1_preds %>%
  group_by(vowel, f13) %>%
  ggplot(aes(f13, duration)) +
  stat_lineribbon() +
  facet_wrap(~vowel) +
  scale_fill_brewer()
```



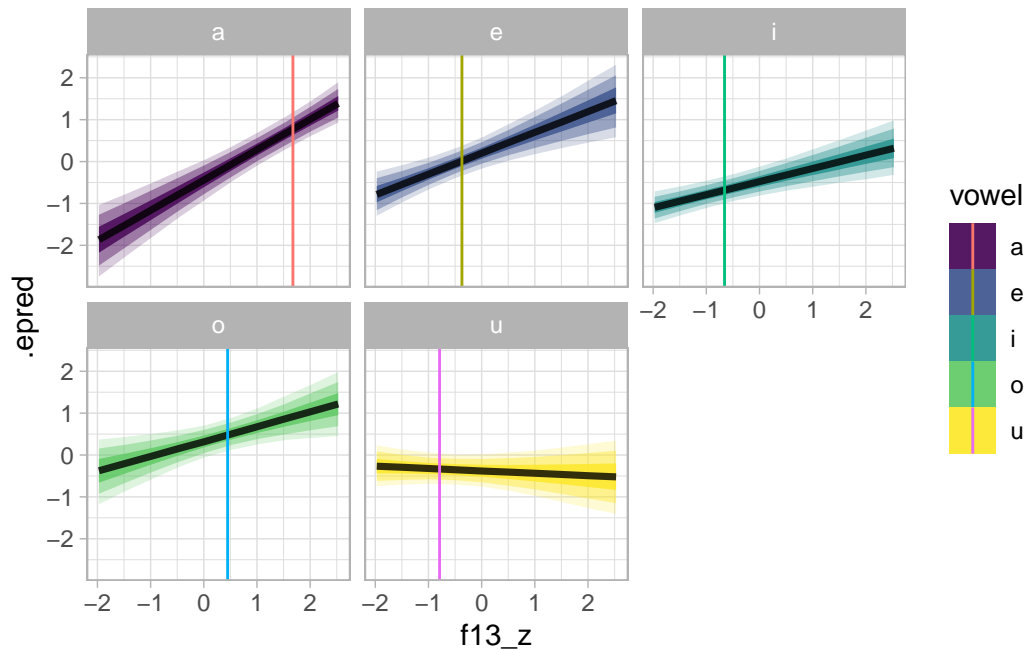
```
bm_1_preds %>%
  group_by(vowel, f13) %>%
  ggplot(aes(f13, duration, fill = vowel)) +
  stat_ribbon(.width = 0.98, alpha = 0.2) +
  stat_ribbon(.width = 0.9, alpha = 0.4) +
  stat_lineribbon(.width = 0.6, alpha = 0.8) +
  geom_vline(data = vmean_f13, aes(xintercept = f13_mean, colour = vowel), linetype = "dashed") +
  geom_rug(data = formants, alpha = 0.1, length = unit(0.015, "npc"), aes(colour = vowel)) +
  facet_wrap(~vowel) +
  labs(
    x = "F1 (Hz)", y = "Duration (ms)"
  ) +
  scale_fill_viridis_d() +
  scale_colour_viridis_d() +
  theme(legend.position = "none")
```



```
ggsave("img/bm1-pred-plot-ms-hz.png", width = 7, height = 5)
```

But let's also plot this in the standardised logged duration scale.

```
bm_1_preds %>%
  group_by(vowel, f13) %>%
  ggplot(aes(f13_z, .epred, fill = vowel)) +
  stat_ribbon(.width = 0.98, alpha = 0.2) +
  stat_ribbon(.width = 0.9, alpha = 0.4) +
  stat_lineribbon(.width = 0.6, alpha = 0.8) +
  geom_vline(data = vmean_f13z, aes(xintercept = f13z_mean, colour = vowel)) +
  facet_wrap(~vowel) +
  scale_fill_viridis_d()
```



### Average predictions and comparisons

```
avg_comparisons(bm_1, variables = "f13_z", conf_level = 0.9)
```

Term	Contrast	Estimate	5.0 %	95.0 %
f13_z	+1	0.357	0.287	0.427

Columns: term, contrast, estimate, conf.low, conf.high

```
avg_comparisons(bm_1, variables = "f13_z", by = "vowel", conf_level = 0.9)
```

Term	Contrast	vowel	Estimate	5.0 %	95.0 %
f13_z	mean(+1)	o	0.3583	0.163	0.5559
f13_z	mean(+1)	a	0.7207	0.575	0.8842
f13_z	mean(+1)	e	0.5006	0.320	0.6735
f13_z	mean(+1)	u	-0.0566	-0.201	0.0906
f13_z	mean(+1)	i	0.3164	0.219	0.4172

Columns: term, contrast, vowel, estimate, conf.low, conf.high, predicted, predicted\_hi, predicted\_lo

```

avg_predictions(bm_1, by = "vowel", conf_level = 0.9) %>%
  as_tibble() %>%
  mutate_if(
    is.numeric, function (x) {exp(x * sd(formants$duration_log) + mean(formants$duration_log))
  })

```

```

# A tibble: 5 x 4
  vowel estimate conf.low conf.high
  <fct>      <dbl>    <dbl>    <dbl>
1 a         99.5     98.4     101.
2 e         79.2     78.3     80.1
3 i         63.7     63.0     64.4
4 o         91.3     90.1     92.5
5 u         70.8     70.0     71.5

```


## Group-level effects

```

bm_1_ranef <- bm_1 %>%
  spread_draws(b_Intercept, r_speaker[subject,var]) %>%
  mutate(condition_mean = b_Intercept + r_speaker)

bm_1_ranef %>%
  filter(var == "f13_z") %>%
  ggplot(aes(y = reorder(subject, r_speaker, median), x = r_speaker)) +
  stat_halfeye()

```



## Non-linear modelling

### F1

```
gam_1 <- bam(
  duration_logz ~
    vowel +
    s(f13_z) +
    s(f13_z, speaker, by = vowel, bs = "fs", m = 1),
  data = formants
)
```

```
summary(gam_1)
```

Family: gaussian

Link function: identity

Formula:

```
duration_logz ~ vowel + s(f13_z) + s(f13_z, speaker, by = vowel,
  bs = "fs", m = 1)
```



Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.01644	0.07064	0.233	0.8160
vowel1	-0.37509	0.18756	-2.000	0.0456 *
vowel2	0.31573	0.14358	2.199	0.0280 *
vowel3	-0.32102	0.12877	-2.493	0.0127 *
vowel4	0.36146	0.16098	2.245	0.0248 *

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
s(f13_z)	4.522	5.439	16.791	<2e-16 ***
s(f13_z,speaker):vowela	25.638	145.000	6.699	<2e-16 ***
s(f13_z,speaker):vowele	24.058	137.000	6.514	<2e-16 ***
s(f13_z,speaker):voweli	27.487	157.000	4.448	<2e-16 ***
s(f13_z,speaker):vowelo	34.267	149.000	6.074	<2e-16 ***
s(f13_z,speaker):vowelu	31.521	151.000	6.066	<2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.712 Deviance explained = 72.7%

fREML = 2663.4 Scale est. = 0.28762 n = 3053

```
vmean <- aggregate(formants$f13_z, list(formants$vowel), mean)
# fs_terms <- c("s(f13_z,speaker)")
fs_terms <- c("s(f13_z,speaker):vowela", "s(f13_z,speaker):vowele", "s(f13_z,speaker):voweli", "s(f13_z,speaker):vowelo", "s(f13_z,speaker):vowelu")

predict_gam(gam_1, exclude_terms = fs_terms, length_out = 100) %>%
  plot(series = "f13_z", comparison = "vowel") +
  geom_vline(data = vmean, aes(xintercept = x, colour = Group.1)) +
  geom_hline(yintercept = 0, linetype = "dashed") +
  facet_wrap(~vowel)
```

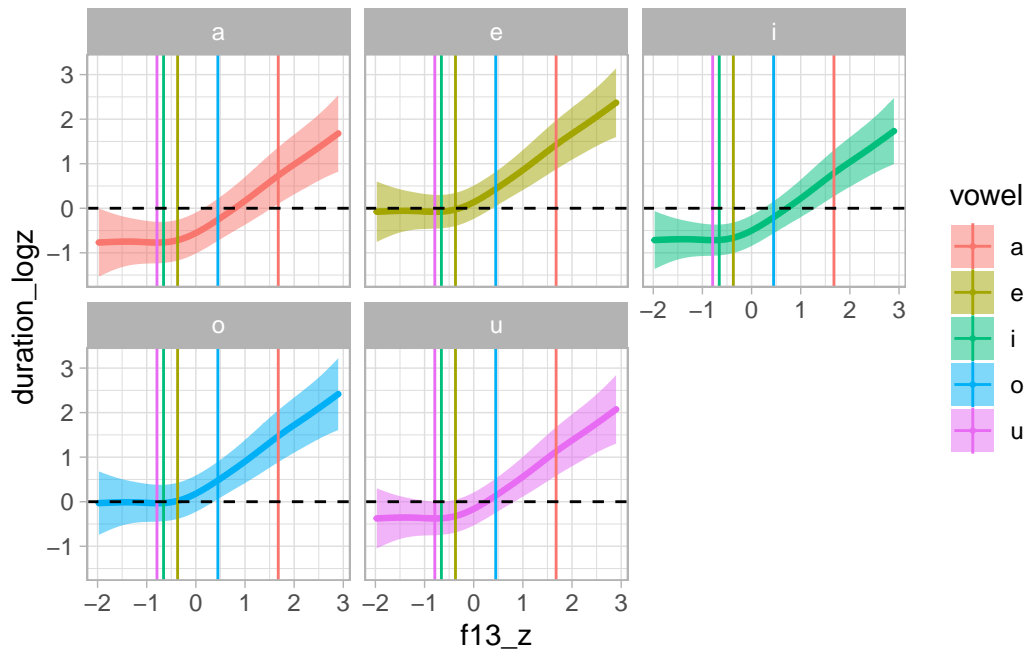
Warning: There was 1 warning in `dplyr::mutate()`.

i In argument: `fit = rowSums(dplyr::across())`.

Caused by warning:

! Using `across()` without supplying `.cols` was deprecated in dplyr 1.1.0.

i Please supply `.cols` instead.



## BRM

```
priors_s <- c(
  prior(normal(0, 1), class = Intercept),
  prior(normal(0, 1), class = b),
  prior(cauchy(0, 0.01), class = sigma),
  prior(cauchy(0, 1), class = sds)
)

bms_1_priors <- brm(
  duration_logz ~
    vowel +
    s(f13_z, k = 5) +
    s(f13_z, speaker, by = vowel, bs = "fs", m = 1, k = 5),
  family = gaussian,
  data = formants,
  prior = priors_s,
  sample_prior = "only",
  cores = 4,
  threads = threading(2),
  backend = "cmdstanr",
  file = "data/cache/bms_1_priors",
)
```

```


    seed = my_seed
)

```

```

conditional_effects(bms_1_priors, "f13_z:vowel")


```



```

conditional_effects(bms_1_priors, "f13_z:vowel", spaghetti = TRUE, ndraws = 100)

```



We specify  $k = 5$  based on the mgcv modelling above. Reducing  $k$  speeds up estimation (because there are less basis functions, hence less parameters to estimate).

The model takes about 4-5 hours to run on 8 cores.

```
bms_1 <- brm(
  duration_logz ~
    vowel +
    s(f13_z, k = 5) +
    s(f13_z, speaker, by = vowel, k = 5, bs = "fs", m = 1),
  family = gaussian,
  data = formants,
  prior = priors_s,
  cores = 4,
  iter = 4000,
  control = list(adapt_delta = 0.9999),
  threads = threading(2),
  backend = "cmdstanr",
  file = "data/cache/bms_1",
  seed = my_seed
)

summary(bms_1, prob = 0.9)
```

Family: gaussian  
 Links: mu = identity; sigma = identity  
 Formula: duration\_logz ~ vowel + s(f13\_z, k = 5) + s(f13\_z, speaker, by = vowel, k = 5, bs =  
 Data: formants (Number of observations: 3053)  
 Draws: 4 chains, each with iter = 4000; warmup = 2000; thin = 1;  
 total post-warmup draws = 8000

Smooth Terms:

	Estimate	Est.Error	l-90% CI	u-90% CI	Rhat	Bulk_ESS	
sds(sf13_z_1)	2.04	1.13	0.96	3.89	1.00	1391	
sds(sf13_zspeakervowela_1)	0.31	0.22	0.03	0.71	1.01	544	
sds(sf13_zspeakervowela_2)	1.88	1.72	0.09	5.25	1.02	486	
sds(sf13_zspeakervoweles_1)	0.34	0.22	0.03	0.73	1.00	374	
sds(sf13_zspeakervoweles_2)	1.69	1.59	0.08	4.91	1.01	493	
sds(sf13_zspeakervoweli_1)	0.33	0.22	0.03	0.72	1.00	475	
sds(sf13_zspeakervoweli_2)	1.82	1.68	0.09	5.16	1.00	591	
sds(sf13_zspeakervowelo_1)	0.32	0.22	0.03	0.73	1.00	414	
sds(sf13_zspeakervowelo_2)	1.75	1.68	0.08	5.15	1.02	455	
sds(sf13_zspeakervowelu_1)	0.32	0.21	0.03	0.70	1.01	460	
sds(sf13_zspeakervowelu_2)	1.66	1.57	0.09	4.88	1.01	654	
	Tail_ESS						
sds(sf13_z_1)	2370						
sds(sf13_zspeakervowela_1)	1266						
sds(sf13_zspeakervowela_2)	1255						
sds(sf13_zspeakervoweles_1)	973						
sds(sf13_zspeakervoweles_2)	1255						
sds(sf13_zspeakervoweli_1)	1100						
sds(sf13_zspeakervoweli_2)	1195						
sds(sf13_zspeakervowelo_1)	1139						
sds(sf13_zspeakervowelo_2)	1055						
sds(sf13_zspeakervowelu_1)	1182						
sds(sf13_zspeakervowelu_2)	1417						

Population-Level Effects:

	Estimate	Est.Error	l-90% CI	u-90% CI	Rhat	Bulk_ESS	Tail_ESS
Intercept	0.06	0.15	-0.20	0.30	1.00	1639	2455
vowel1	-0.18	0.10	-0.35	-0.01	1.01	1300	2422
vowel2	0.29	0.04	0.22	0.36	1.01	1556	3821
vowel3	-0.40	0.04	-0.47	-0.34	1.01	1537	2909
vowel4	0.34	0.04	0.27	0.40	1.00	2819	3944
sf13_z_1	0.90	0.69	-0.23	2.04	1.00	2538	3869

Family Specific Parameters:

	Estimate	Est.Error	l-90% CI	u-90% CI	Rhat	Bulk_ESS	Tail_ESS
sigma	0.55	0.01	0.54	0.56	1.00	3441	4293

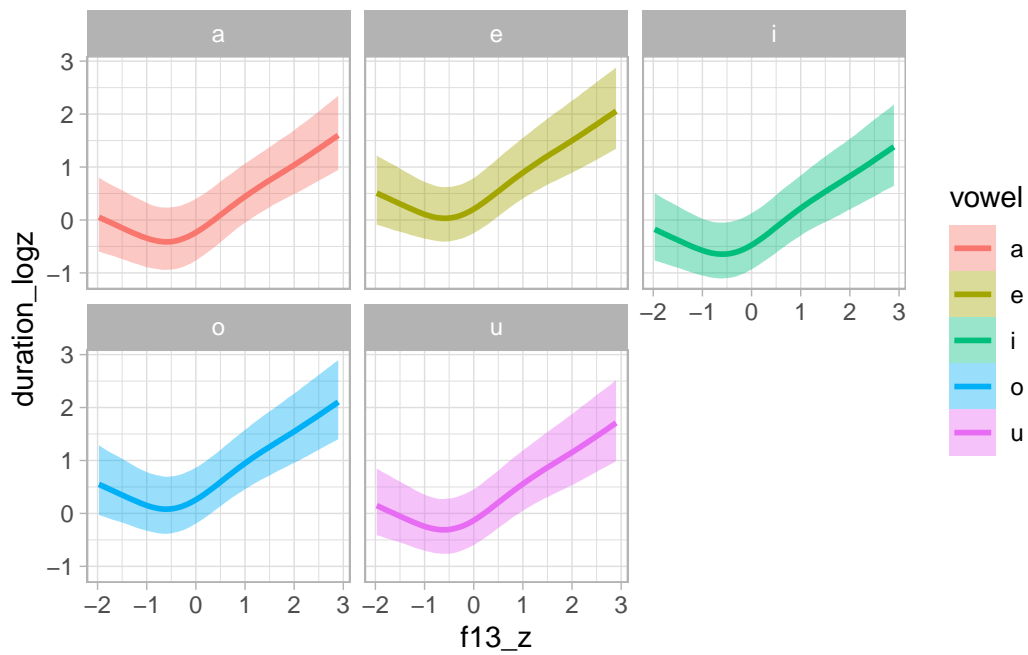
Draws were sampled using `sample(hmc)`. For each parameter, `Bulk_ESS` and `Tail_ESS` are effective sample size measures, and `Rhat` is the potential scale reduction factor on split chains (at convergence, `Rhat = 1`).

```
bms_1 %>%
  as_draws_df() %>%
  select(b_vowel1:b_vowel4) %>%
  pivot_longer(b_vowel1:b_vowel4) %>%
  group_by(name) %>%
  summarise(
    cri95 = list(round(quantile2(value, probs = c(0.025, 0.975)), 2)),
    cri90 = list(round(quantile2(value, probs = c(0.05, 0.95)), 2)),
    cri80 = list(round(quantile2(value, probs = c(0.1, 0.9)), 2)),
    cri60 = list(round(quantile2(value, probs = c(0.2, 0.8)), 2))
  ) %>%
  knitr::kable(format = "latex") %>% cat(sep = "\n")
```


Warning: Dropping 'draws\_df' class as required metadata was removed.

```
\begin{tabular}{l|l|l|l|l|l|l}
\hline
name & cri95 & cri90 & cri80 & cri60 \\
\hline
b\_vowel1 & -0.38, 0.02 & -0.35, -0.01 & -0.31, -0.05 & -0.26, -0.09 \\
\hline
b\_vowel2 & 0.21, 0.37 & 0.22, 0.36 & 0.24, 0.34 & 0.26, 0.32 \\
\hline
b\_vowel3 & -0.48, -0.32 & -0.47, -0.34 & -0.45, -0.35 & -0.43, -0.37 \\
\hline
b\_vowel4 & 0.26, 0.41 & 0.27, 0.40 & 0.29, 0.39 & 0.30, 0.37 \\
\hline
\end{tabular}
```

```
plot(conditional_effects(bms_1, "f13_z:vowel"), plot = FALSE)[[1]] + facet_wrap(~vowel)
```



```
plot(conditional_effects(bms_1, "f13_z:vowel", spaghetti = TRUE, ndraws = 100), plot = FALSE)
```



Let's plot on the original scale.

```


bms_1_grid <- expand_grid(
  vowel = levels(formants$vowel),
  f13_z = seq_minmax(formants$f13_z, 0.25),
  speaker = NA
)

bms_1_preds <- epred_draws(bms_1, newdata = bms_1_grid, re_formula = NA) %>%
  mutate(
    duration_log = .epred * sd(formants$duration_log) + mean(formants$duration_log),
    duration = exp(duration_log),
    f13 = f13_z * sd(formants$f13) + mean(formants$f13)
  )

bms_1_preds %>%
  group_by(vowel, f13) %>%
  ggplot(aes(f13, duration, fill = vowel)) +
  stat_ribbon(.width = 0.98, alpha = 0.2) +
  stat_ribbon(.width = 0.9, alpha = 0.4) +
  stat_lineribbon(.width = 0.6, alpha = 0.8) +
  geom_vline(data = vmean_f13, aes(xintercept = f13_mean, colour = vowel), linetype = "dashed") +
  geom_rug(data = formants, alpha = 0.1, length = unit(0.015, "npc"), aes(colour = vowel)) +
  facet_wrap(~vowel) +
  labs(
    x = "F1 (Hz)", y = "Duration (ms)"
  ) +
  scale_fill_viridis_d() +
  scale_colour_viridis_d() +
  theme(legend.position = "none")

```





```
ggsave("img/bms1-pred-plot-ms-hz.png", width = 7, height = 5)
```

## F1 and F2

```
gam_2 <- bam(
  duration_logz ~
    vowel +
    s(f13_z, f23_z) +
    s(f13_z, f23_z, speaker, bs = "fs", m = 1),
  data = formants
)
```

```
summary(gam_2)
```

Family: gaussian

Link function: identity

Formula:

duration\_logz ~ vowel + s(f13\_z, f23\_z) + s(f13\_z, f23\_z, speaker,

```
bs = "fs", m = 1)
```

Parametric coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.04815	0.14310	0.336	0.737
vowel1	-0.14641	0.09984	-1.466	0.143
vowel2	0.40133	0.05185	7.741	1.35e-14 ***
vowel3	-0.33404	0.05005	-6.675	2.95e-11 ***
vowel4	0.25247	0.04908	5.144	2.87e-07 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:

	edf	Ref.df	F	p-value
s(f13_z,f23_z)	13.1	17.04	6.503	<2e-16 ***
s(f13_z,f23_z,speaker)	101.9	567.00	6.987	<2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.708 Deviance explained = 71.9%

fREML = 2606.4 Scale est. = 0.29247 n = 3053

```
gam_2_preds <- predict_gam(gam_2, length_out = 50, exclude_terms = "s(f13_z,f23_z,speaker)
```

```
vmeans <- formants %>%  
  group_by(vowel) %>%  
  summarise(  
    f13_z = mean(f13_z), f23_z = mean(f23_z)  
  )
```

```
gam_2_preds %>%  
  ggplot(aes(f23_z, f13_z)) +  
  geom_raster(aes(fill = duration_logz), interpolate = TRUE) +  
  geom_contour(aes(z = duration_logz), bins = 40, colour = "white", linewidth = 0.05) +  
  geom_label(data = vmeans, aes(label = vowel), size = 5) +  
  scale_x_reverse(position = "top") +  
  scale_y_reverse(position = "right") +  
  scale_fill_distiller(palette = "BuPu")
```

